

CS 250: Computer Architecture

Final Exam

Spring 2025

Benjamin Lobos Lertpunyaroj

May 8th, 10:30_{AM} – 12:30_{PM}

Exam contents and details for referencing

- Final exam is held in Fowler Hall on May 8th (Thursday), from 10:30 _{AM} to 12:30 _{PM}.
- Previous cumulative book chapters
 - Chapter 1 sections 1, 2, and 3.
 - Chapter 2, sections 1, 2, 3, 4, 5, 6, and 7.
 - Chapter 3, sections 1, 2, and 5.
 - Chapter 4, sections 1, 2, 3, 4, 5, 6, 7, and 8.
 - Chapter 5, sections 1, 2, 3, 4, 7, and 8.
 - Chapter 8 (Appendix A), sections 1, 2, 3 (but not PLAs or ROMs), 5, 7 (lightly), and 8.
- All lecture notes and lecture slides.
- All labs (1 - 11).

Appendix A

Logic & Gates

An *asserted* signal is logically true, the *deasserted* is the opposite.

Two types of logic systems

Combinational logic: No memory in components, hence same output given same input.

Sequential logic: Memory in components, hence output depends on input and current memory state.

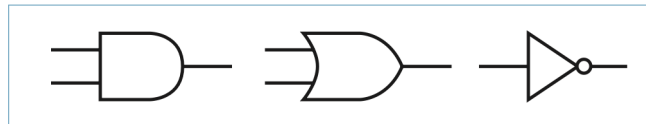


Figure 1: AND gate, OR gate, and inverter

The gates can be combined to form different forms of logic. An example of this is $\overline{\overline{A} + \overline{B}}$ which is equivalent to $A \cdot B$ by De Morgan's law, seen in Figure 2.

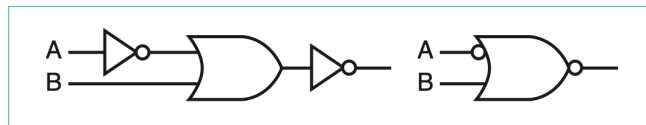


Figure 2: Logic gate implementation of example formula

Decoders & Multiplexors

A **decoder** is a logic block that has an n -bit input and 2^n outputs, where there is one unique true bit as output from a unique set of bytes of input.

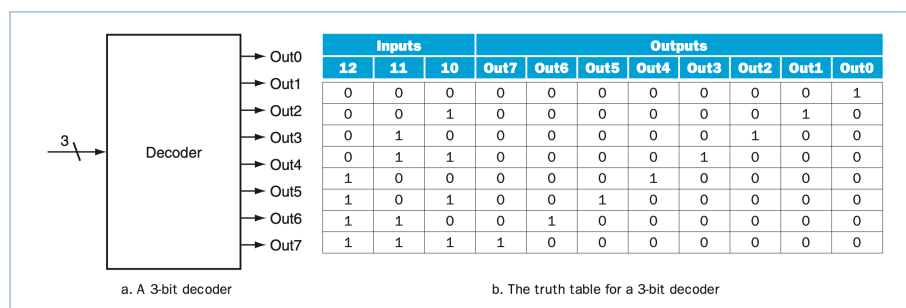


Figure 3: 3-bit input decoder that generates $2^3 = 8$ different outputs (Out0 – Out7)

$$2^n \text{ outputs} \therefore \log_2(\text{output}) = \text{input bits}$$

Encoders are the other way around.

Multiplexors have a selector input (or control value), that will determine which inputs will become outputs.

In the case of the two-input MUX, its representation is the following, $C = (A \cdot \bar{S}) + (B \cdot S)$, using n (data inputs) AND gates, and one OR gate.

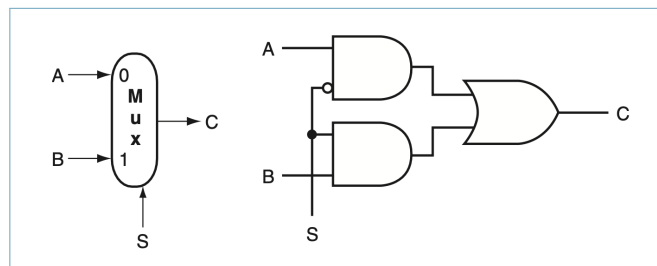


Figure 4: *Two-input multiplexor that generates one output depending on the selector input S*

$$n \text{ (data inputs)} \therefore \log_2 n = S \text{ selector bits required to represent all inputs}$$

Often times a decoder generates n bits for a MUX, to be used as a selector signal.

Buses

A collection of data lines that is treated as a single logical signal.

When showing a logic unit whose inputs and outputs are buses, the unit must be replicated a sufficient number of times to accommodate the width of the input.

You can use multiplexors to select between two buses, requiring n inputs to represent n -bit buses.

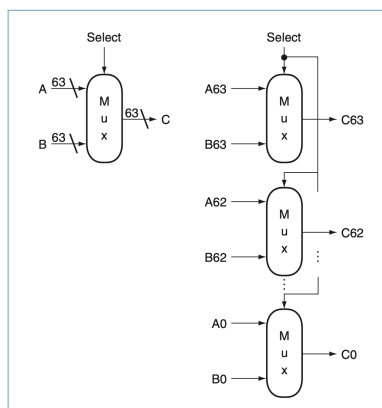


Figure 5: *1-bit multiplexors replicated 64 times to represent two 64-bit buses*

ALUs

Operation done by the ALU

Logic operations: AND and OR gate operations, with NOR being available through an inversion of both input signals with AInvert and BInvert control signals.

Arithmetic operations: Addition and subtraction through the full adder, and BInvert control signal on one input for determining the type of operation.

The LEGv8 word is 64 bits wide, as such a 64 bit wide ALU is required (64 1-bit ALUs).

In its simplest form, a 1-bit logical unit for AND and OR operations simply requires a multiplexor and a one bit control signal to select between the two operations ($2^1 = 2$).

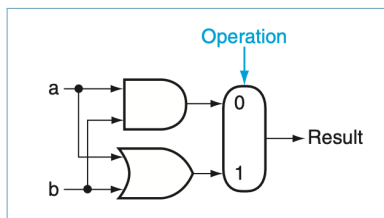


Figure 6: 1-bit logical unit for AND and OR operations

Implementing addition requires two input operands, one output, a CarryIn bit carried from the less-significant bits of the operation (i.e. another 1-bit logical unit), and a CarryOut bit to be carried forward to the next more significant bit.

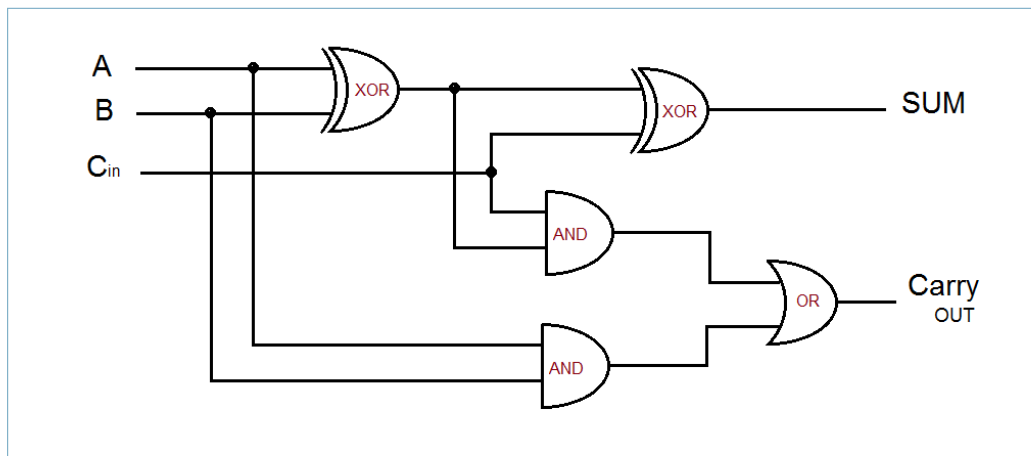


Figure 7: Full adder that performs mod 2 addition

The combination of the adder and the logic gates, coupled with a multiplexor with a control signal to determine the operation makes a complete 1-bit ALU, which can be seen in Figure 8.

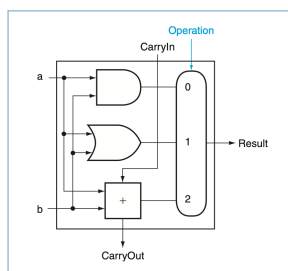


Figure 8: 1-bit alu with logical operations and addition

For expanding to a 64-bit ALU, the adders have to set up a ripple carry from the least to the most significant bit.

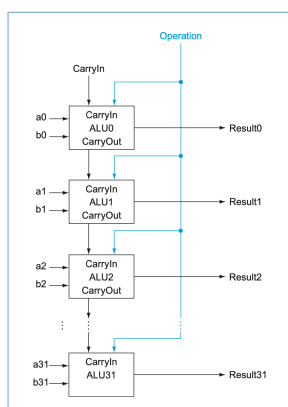


Figure 9: Ripple carry implemented for a 64-bit ALU

By inverting the second input (**BInvert** = 1, seen in Figure 10) and setting **CarryIn** to 1 in the least significant bit of the ALU, we get two's complement subtraction of **b** from **a**.

To implement a NOR function, existing components can be combined, $\overline{(a + b)} = \bar{a} \cdot \bar{b}$ (DeMorgan's theorem), which means we need an AND and two inverters for both **a** and **b**, seen in Figure 10

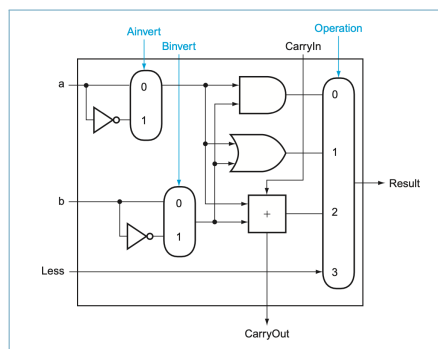


Figure 10: 1-bit ALU that performs subtraction, and NOR operations

On a 64-bit ALU we can use a zero flag to help with conditional branch instructions in LEGv8 (e.g.

CBZ), as they receive to inputs and require to test if the subtraction has a zero.

The following represents this with an inversion of an OR tree on all results from the subtraction considering a 64-bit subtraction, fully represented in Figure 11.

$$Zero = \overline{(R_0 + R_1 + R_2 + \dots + R_{63})}$$

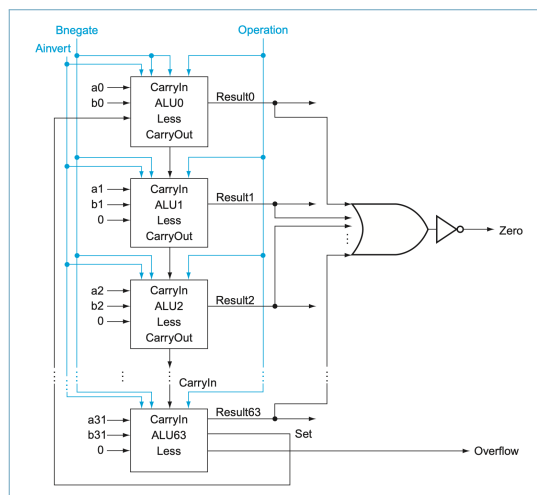


Figure 11: 64-bit ALU OR tree and an inverter for determining the Zero flag

For a generalized symbol of the ALU, Figure 12, where ALU operation is the control signal of the MUX that determines the type of operation.

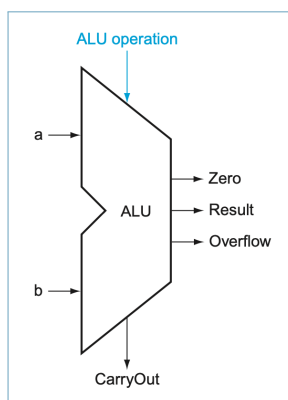


Figure 12: General symbol for an ALU or an adder

Clocks

Clocking methodology semantics

Edge triggered clocking: State changes occur on a clock edge.

Synchronous system: Type of memory system where data is read only when a clock signal indicates stability (i.e. non-changing value).

A combinational logic block, receives an input and then generates an output for a state element which is updated on a clock edge.

An edge-triggered methodology allows a state element to be read and written in the same clock cycle without creating a race condition.

For this to work, the clock cycle must be long enough for the state element to have received a stable input before the next active clock edge.

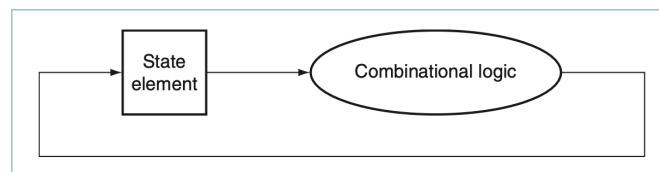


Figure 13: *Edge-triggered state element to be read and written to in one active clock edge*

One such state element is the register file.

Flip-flops & Latches

Types of clocked memory elements

Flip-flops: Edge-triggered element that changes the stored state only at a clock edge.

Latches: Level-sensitive element that changes the stored state at any time the clock is asserted.

Flip-flops are build upon latches and are going to be used in edge-triggered systems.

A **D flip-flop** or **D latch** is used for storing the value of one data input signal, in the internal memory, at the clock edge.

To implement a **D latch**, it requires two inputs, the data to be stored D , and the clock signal C , producing two outputs, the value of the internal state Q , and its complement \bar{Q} .

The implementation has cross-coupled **NOR** gates that store the state value unless C is asserted, in which case D replaces the value of Q and is stored.

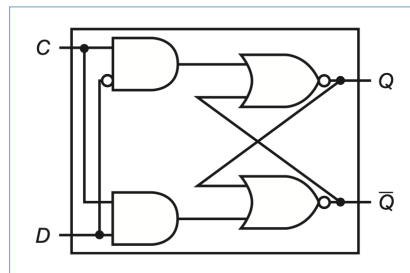


Figure 14: *D latch, composed of crossed NOR gates and a SR latch*

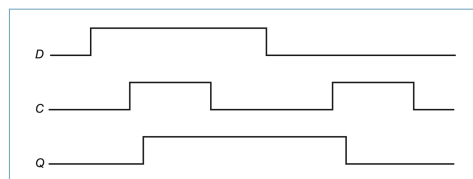


Figure 15: *Progression of a D latch, assuming output is initially deasserted*

To implement a **D flip-flop**, with a falling-edge trigger, we can use two D latches, master and slave. Master sets input D when C is asserted. When C falls, master is closed, but slave is open and gets its input from master's Q.

In this sense, the rising-edge represents when the master takes in the D value, and the falling-edge represents when the slave takes in the master's D producing the final Q.

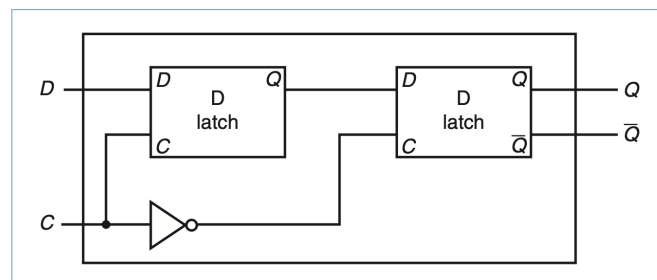


Figure 16: *D flip-flop with a falling-edge trigger made from two D latches, master and slave*




Figure 17: *Progression of a D flip-flop with a falling-edge trigger, where output is initially deasserted*

The minimum time D must retain a valid input is the setup time plus the hold time (after edge).

Register files

A **register file** consists of a bunch of **registers** that can be read and written to, and a **WriteReg** control signal (clock).

For writing it requires the control signal, the number of the register to write to (**Write register**), and the data to write (**Write data**).

For reading it requires the numbers of the registers to read from (**Read register number 1 & 2**), and it outputs the read contents from two registers (**Read data 1 & 2**).

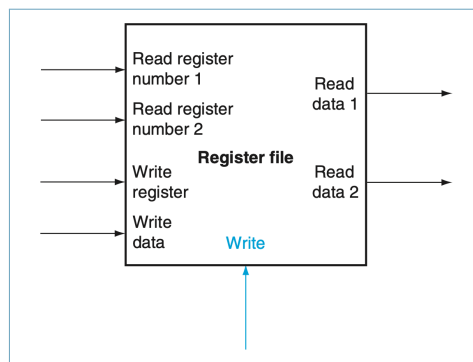


Figure 18: *Register file with two read ports and one write port*

The implementation for the write port consists of a decoder that will select one of the $n - 1$ registers that will be ANDed with the **WriteReg** signal to act as the **C** input for the registers (**D** flip-flops).

The **D** input for every register is the **Write data** input from the reg. file.

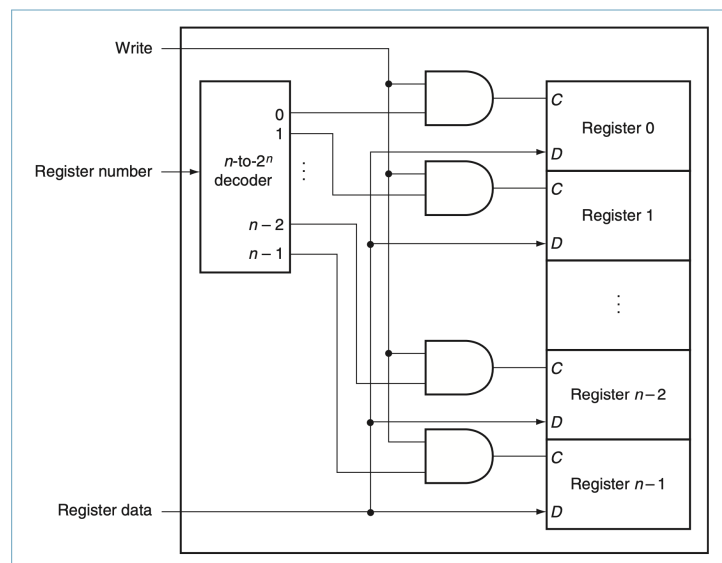


Figure 19: *Write implementation in the register file*

The implementation for the two read ports consists of using the stored state of the registers (Q output), as inputs for two different MUXes that use the Read register number 1 & 2 reg. file inputs as control signals to output the information of the two registers specified.

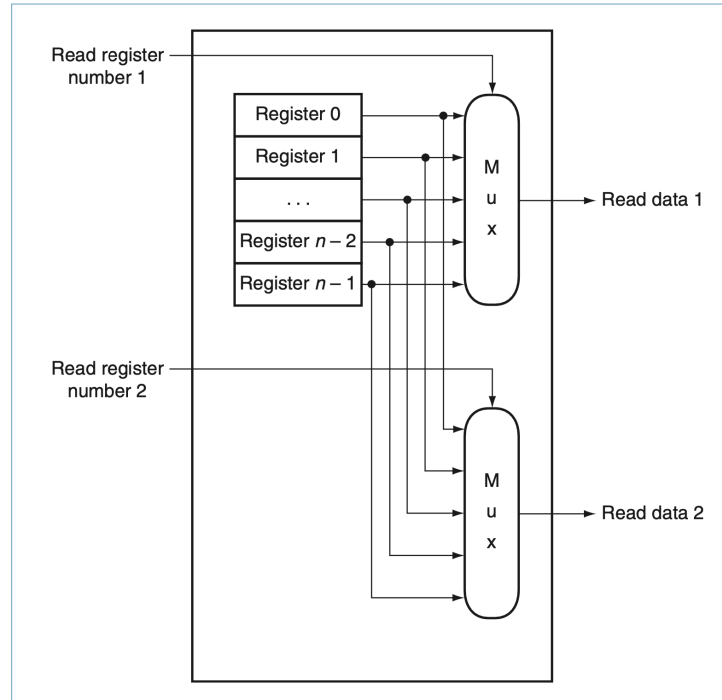


Figure 20: *Read implementation in the register file*